


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error as mse, r2_score
from sklearn.metrics import accuracy_score
import plotly.express as px
```

```
df = pd.read_csv("/content/housing.csv")
df
```



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_val
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	45260
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	35850
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	35210
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	34130
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	34220
...
20635	-121.09	39.48	25.0	1665.0	374.0	845.0	330.0	1.5603	7810
20636	-121.21	39.49	18.0	697.0	150.0	356.0	114.0	2.5568	7710
20637	-121.22	39.43	17.0	2254.0	485.0	1007.0	433.0	1.7000	9230
20638	-121.32	39.43	18.0	1860.0	409.0	741.0	349.0	1.8672	8470
20639	-121.24	39.37	16.0	2785.0	616.0	1387.0	530.0	2.3886	8940


20640 rows × 10 columns

Next steps:

[Generate code with df](#)

 [View recommended plots](#)

```
df.isnull().sum()
```




longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	207
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0
dtype:	int64

```
val = df['total_bedrooms'].median()
val
```



435.0


```
df['total_bedrooms'].dropna()
```



0	129.0
1	1106.0
2	190.0
3	235.0
4	280.0
...	...
20635	374.0
20636	150.0
20637	485.0
20638	409.0
20639	616.0

Name: total_bedrooms, Length: 20433, dtype: float64

```
df.isnull().sum()
```



longitude	0
latitude	0

```
housing_median_age    0
total_rooms            0
total_bedrooms        207
population            0
households            0
median_income         0
median_house_value    0
ocean_proximity       0
dtype: int64
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   longitude            20640 non-null  float64
1   latitude             20640 non-null  float64
2   housing_median_age   20640 non-null  float64
3   total_rooms          20640 non-null  float64
4   total_bedrooms       20433 non-null  float64
5   population           20640 non-null  float64
6   households           20640 non-null  float64
7   median_income        20640 non-null  float64
8   median_house_value   20640 non-null  float64
9   ocean_proximity      20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

df['ocean_proximity'].unique()

```
array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
      dtype=object)
```

```
df['ocean_proximity'] = df['ocean_proximity'].map({'NEAR BAY': 0, '<1H OCEAN': 1, 'INLAND': 2, 'NEAR OCEAN': 3, 'ISLAND': 4})
df['ocean_proximity']
```

```
0      0
1      0
2      0
3      0
4      0
..
20635   2
20636   2
20637   2
20638   2
20639   2
Name: ocean_proximity, Length: 20640, dtype: int64
```


EDA

df.head()

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	322.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0

Next steps: [Generate code with df](#) [View recommended plots](#)

```
fig = px.histogram(df, x="median_house_value")
fig.show()
```




```
corr = df.corr(method = 'pearson')['median_house_value']  
corr
```




```
longitude      -0.045967  
latitude       -0.144160  
housing_median_age  0.105623  
total_rooms    0.134153  
total_bedrooms 0.049686  
population     -0.024650  
households     0.065843  
median_income  0.688075  
median_house_value 1.000000  
ocean_proximity -0.210600  
Name: median_house_value, dtype: float64
```

```
heatmap = sns.heatmap(df.corr(),annot=True)
```



```
data = df.head(5000)  
x = data['median_income']  
y = data['median_house_value']  
px.scatter(data,x,y)
```




```
np.linspace(y.min(), y.max(), 11)
```



```
array([ 14999. ,  63499.2, 111999.4, 160499.6, 208999.8, 257500. ,  
       306000.2, 354500.4, 403000.6, 451500.8, 500001. ])
```

```
plt.boxplot(data['median_house_value'])  
plt.show()
```

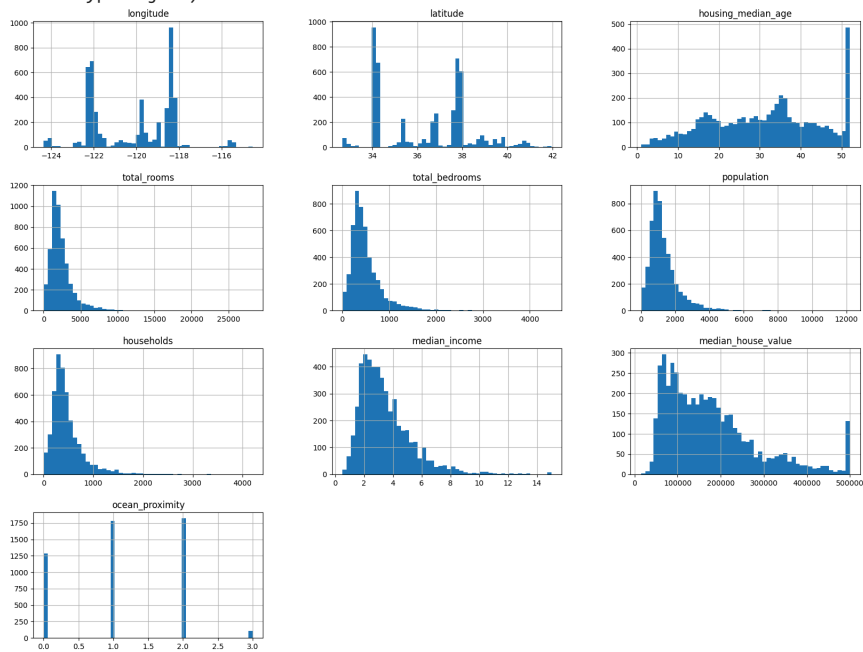


```
data.hist(bins=50,figsize=(20,15))
```

```

array([[<Axes: title={'center': 'longitude'}>,
        <Axes: title={'center': 'latitude'}>,
        <Axes: title={'center': 'housing_median_age'}>],
       [[<Axes: title={'center': 'total_rooms'}>,
        <Axes: title={'center': 'total_bedrooms'}>,
        <Axes: title={'center': 'population'}>],
       [[<Axes: title={'center': 'households'}>,
        <Axes: title={'center': 'median_income'}>,
        <Axes: title={'center': 'median_house_value'}>],
       [[<Axes: title={'center': 'ocean_proximity'}>, <Axes: >, <Axes: >]],
       dtype=object)

```



```

(df['median_house_value'] > 420000)

```

```

0      True
1     False
2     False
3     False
4     False
...
20635  False
20636  False
20637  False
20638  False
20639  False
Name: median_house_value, Length: 20640, dtype: bool

```

```

(df['median_house_value'] > 420000).sum()

```

```

1540

```

```
df.columns
housing_Df = df[['longitude', 'latitude', 'housing_median_age', 'total_rooms',
                 'total_bedrooms', 'population', 'households', 'median_income',
                 'median_house_value', 'ocean_proximity']]
housing_Df.head()
housing_Df.dropna(inplace=True)
housing_Df.isnull().sum()
```

```
longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms  0
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
train_pd, test_pd, val_pd = housing_Df[:18000], housing_Df[18000:19217], housing_Df[19215:]
len(train_pd), len(test_pd), len(val_pd)
```

```
(18000, 1217, 1218)
```

```
X_train, y_train = train_pd.drop('median_house_value', axis=1), train_pd.to_numpy()[:-1]
X_train.head(10)
```

```
longitude latitude housing_median_age total_rooms total_bedrooms population
0 -122.23 37.88 41.0 880.0 129.0 322.0
1 -122.22 37.86 21.0 7099.0 1106.0 2401.0
2 -122.24 37.85 52.0 1467.0 190.0 496.0
3 -122.25 37.85 52.0 1274.0 235.0 558.0
4 -122.25 37.85 52.0 1627.0 280.0 565.0
5 -122.25 37.85 52.0 919.0 213.0 413.0
6 -122.25 37.84 52.0 2535.0 489.0 1094.0
7 -122.25 37.84 52.0 3104.0 687.0 1157.0
8 -122.26 37.84 42.0 2555.0 665.0 1206.0
9 -122.25 37.84 52.0 3549.0 707.0 1551.0
```

Next steps: [Generate code with X_train](#) [View recommended plots](#)

```
X_val, y_val = val_pd.to_numpy()[:-1], val_pd.to_numpy()[:-1]
X_test, y_test = test_pd.to_numpy()[:-1], test_pd.to_numpy()[:-1]
```

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape, X_val.shape, y_val.shape
```

```
((18000, 9), (18000,), (1217, 9), (1217,), (1218, 9), (1218,))
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)
X_val = scaler.fit_transform(X_val)
```


```
pd.DataFrame(X_train)
```



	0	1	2	3	4	5	6	
0	-1.453822	1.204250	0.935040	-0.795703	-0.964371	-0.970876	-0.972988	2.32936
1	-1.448767	1.194570	-0.642170	2.018961	1.315597	0.843015	1.637958	2.31688
2	-1.458876	1.189729	1.802506	-0.530032	-0.822019	-0.819064	-0.841409	1.76947
3	-1.463931	1.189729	1.802506	-0.617382	-0.717005	-0.764970	-0.733049	0.92303
4	-1.463931	1.189729	1.802506	-0.457617	-0.611991	-0.758863	-0.629850	-0.01915
...
17995	-1.352724	0.947718	-0.799891	0.530841	1.597968	0.432076	1.676658	-0.19260
17996	-1.352724	0.957398	-1.036472	0.345732	1.518624	0.038586	1.392859	-0.15553
17997	-1.347669	0.952558	-0.642170	-0.075631	0.314465	0.044694	0.479544	0.01104
17998	-1.347669	0.952558	-0.326728	-0.255309	-0.362291	-0.141145	-0.309932	0.46045
17999	-1.342614	0.952558	-1.036472	-0.693870	-0.439302	-0.816447	-0.462151	0.16955

18000 rows x 9 columns

```
pd.DataFrame(X_test).hist(bins=50, figsize=(20,15))  
plt.show()
```



```
X_train.shape, X_test.shape, X_val.shape
```



```
((18000, 9), (1217, 9), (1218, 9))
```

Linear Regression

```
# Preprocessing - scaling the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)
# Train the model
```

```
y_train_pred = lm.predict(X_train_scaled)
y_val_pred = lm.predict(X_val_scaled)
```

```
mse_train = mse(y_train, y_train_pred)
rmse_train = mse(y_train, y_train_pred, squared=False)
# Calculate MSE and RMSE for validation set
mse_val = mse(y_val, y_val_pred)
rmse_val = mse(y_val, y_val_pred, squared=False)
# Calculate R2 score for training set
r2_train = r2_score(y_train, y_train_pred)
# Calculate R2 score for validation set
r2_val = r2_score(y_val, y_val_pred)
```

```
print(f'Training MSE: {mse_train}')
print(f'Training RMSE: {rmse_train}')
print(f'Training R2: {r2_train}')
```

```
➞ Training MSE: 3.543895986947175e-31
   Training RMSE: 5.953063066142652e-16
   Training R2: 1.0
```

```
print(f'Validation MSE: {mse_val}')
print(f'Validation RMSE: {rmse_val}')
print(f'Validation R2: {r2_val}')
```

```
➞ Validation MSE: 1.4624634571995185
   Validation RMSE: 1.2093235535618738
   Validation R2: -3.725684670880182
```

```
print('Some predictions on the validation set:', y_val_pred[:5])
```

```
➞ Some predictions on the validation set: [2.50073119 1.5260041 1.6479579 1.66873522 1.61001672]
```

```
plt.figure(figsize=(10, 5))
# Plot actual values
plt.scatter(range(len(y_val)), y_val, color='blue', alpha=0.5, label='Actual')
# Plot predicted values
plt.scatter(range(len(y_val)), y_val_pred, color='red', alpha=0.5, label='Predicted')
plt.xlabel('House Price')
plt.ylabel('Value')
plt.title('Actual vs Predicted Values')
```