

Chinmay Rajan Mhatre

22102B2001

CMPN B

https://github.com/chinmay0910/ML-Lab---VIT/blob/main/EXP%208/ML_LAB_8.ipynb


✓ Importing required libraries

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sb
4 import matplotlib.pyplot as pt
5
6 from sklearn.cluster import DBSCAN
7 from sklearn.preprocessing import StandardScaler
8 from sklearn import metrics
9
10 import warnings
11 warnings.simplefilter("ignore")
```


✓ i) Exploratory Data Analysis

```
1 # Loading dataset or dataframe
2 segment=pd.read_csv("/content/Mall_Customers.csv")
```



```
1 # Looking for shape of dataframe
2 segment.shape
```

 (200, 5)


```
1 # Viewing columns
2 segment.columns
```

 Index(['CustomerID', 'Gender', 'Age', 'Annual Income (k\$)',
 'Spending Score (1-100)'],
 dtype='object')

```
1 # Head of the dataframe
2 segment.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40





Next steps:

[Generate code with segment](#)


[View recommended plots](#)

[New interactive sheet](#)

```
1 # Tail of the dataframe
2
3 segment.tail()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
195	196	Female	35	120	79
196	197	Female	45	126	28
197	198	Male	32	126	74
198	199	Male	32	137	18
199	200	Male	30	137	83



```
1 # Datatypes involved..
2 segment.dtypes.value_counts()
```

```
↗
count
int64    4
object    1

dtype: int64
```

```
1 # Information about Dataframe
2 segment.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerID            200 non-null   int64
1   Gender                 200 non-null   object
2   Age                    200 non-null   int64
3   Annual Income (k$)     200 non-null   int64
4   Spending Score (1-100) 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
1 # Removing white spaces and remaining columns
2 segment.columns=segment.columns.str.replace(" ", "")
3 segment.columns
```

```
↗ Index(['CustomerID', 'Gender', 'Age', 'AnnualIncome(k$)',
        'SpendingScore(1-100)'],
        dtype='object')
```

```
1 # renaming columns
2 segment.columns=segment.rename(columns={'AnnualIncome(k$)': 'AnnualIncome',
3                                         'SpendingScore(1-100)': 'SpendingScore', "Genre": "Gender"}).columns
4 segment.columns
```

```
↗ Index(['CustomerID', 'Gender', 'Age', 'AnnualIncome', 'SpendingScore'], dtype='object')
```

```
1 # Viewing int columns
2 int_col=segment.select_dtypes(include="int64").columns.tolist()
3 int_col
```

```
↗ ['CustomerID', 'Age', 'AnnualIncome', 'SpendingScore']
```

```
1 # Viewing categorical columns
2 cat_col=segment.select_dtypes(include="O").columns.tolist()
3 cat_col
```

```
↗ ['Gender']
```

```
1 # Drop the id column
2 copy_segment=segment.copy()
3 segment.drop("CustomerID",axis=1,inplace=True)
```

```
1 # Summary statistics
2 segment.describe()
```

```
↗
```

	Age	AnnualIncome	SpendingScore
count	200.000000	200.000000	200.000000
mean	38.850000	60.560000	50.200000
std	13.969007	26.264721	25.823522
min	18.000000	15.000000	1.000000
25%	28.750000	41.500000	34.750000
50%	36.000000	61.500000	50.000000
75%	49.000000	78.000000	73.000000
max	70.000000	137.000000	99.000000

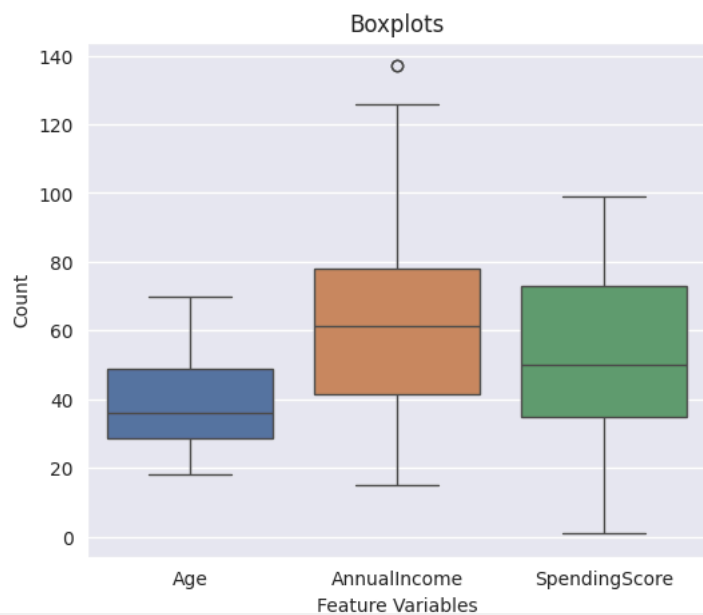
1 Start coding or [generate](#) with AI.

```

1 # looking for outliers through boxplot
2 sb.set({"figure.figsize":(6,5)})
3 sb.boxplot(segment)
4 pt.title("Boxplots")
5 pt.xlabel("Feature Variables")
6 pt.ylabel("Count")
7

```

Text(0, 0.5, 'Count')



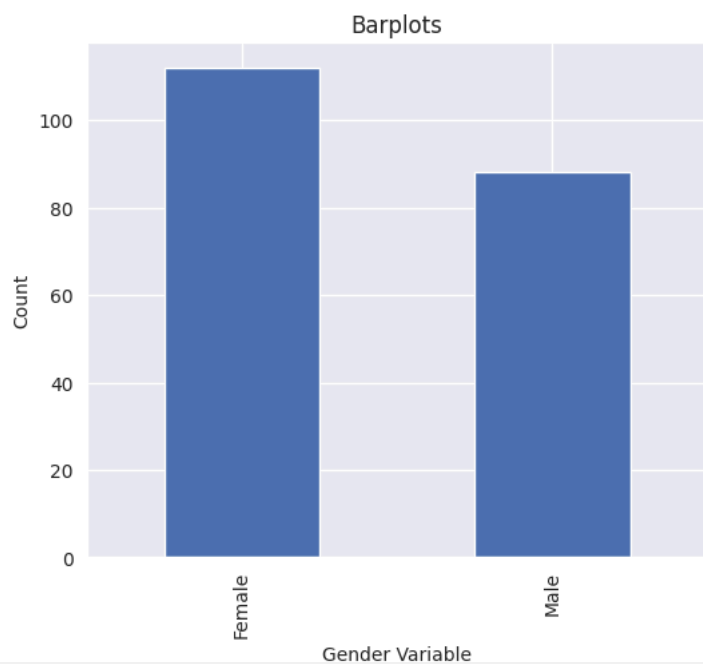
since,we have one outlier in annual income,we neglect it as it may be the person with higher annual income.

```

1 # Plotting gender for count
2 segment.Gender.value_counts().plot(kind="bar")
3 pt.title("Barplots")
4 pt.xlabel("Gender Variable")
5 pt.ylabel("Count")
6

```

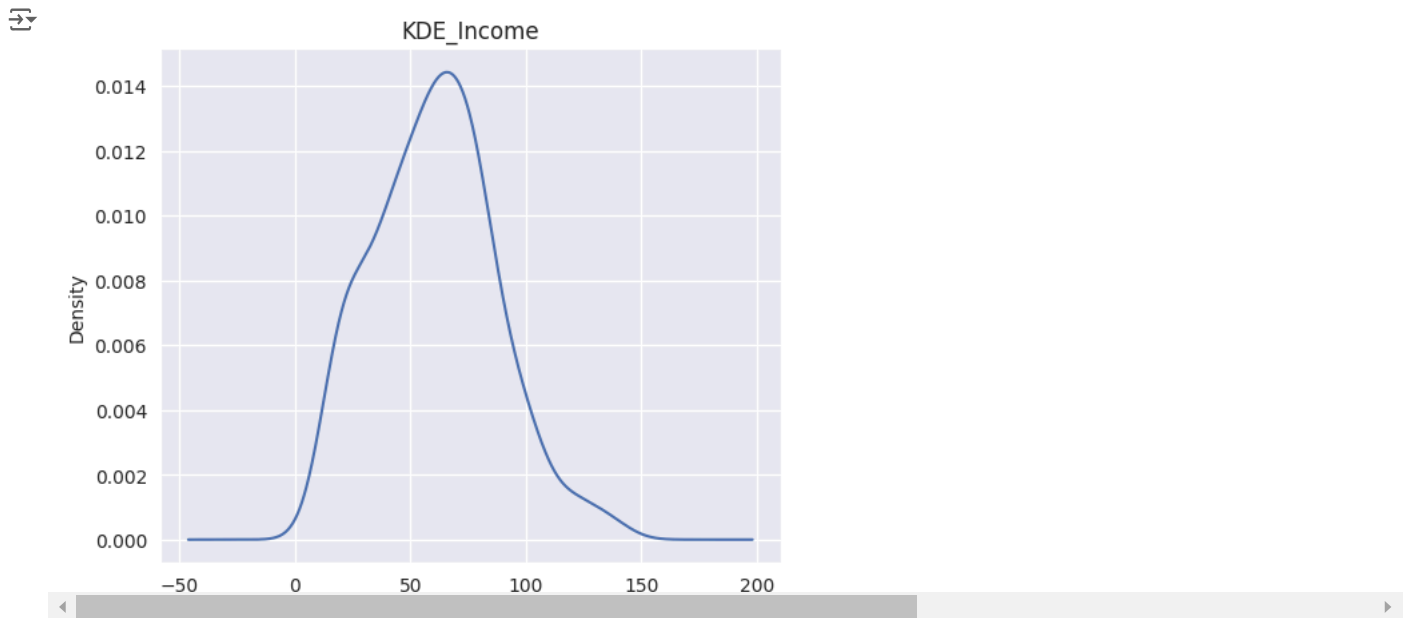
Text(0, 0.5, 'Count')



```

1 # KDE plot for Annual income
2 segment.AnnualIncome.plot(kind="kde")
3 pt.title('KDE_Income')
4 pt.show()

```



```
1 # Scatter plot for the AnnualIncome & SpendingScore
2 sb.set({"figure.figsize":(10,5)})
3 sb.scatterplot(data=segment,x='AnnualIncome',y='SpendingScore',hue="Gender")
4 pt.title('Income vs Score')
5 pt.show()
```



~: We can see that most of the customer data points lies at annual income(40-70) and spending score (40-60).

```
1 # Scatter plot for the Age & SpendingScore
2 sb.scatterplot(data=segment,x='Age',y='SpendingScore',hue="Gender")
3 pt.title('Score vs Age')
4 pt.show()
```



~: We interpret that the age (40-60) of having spending score around (20-60), the age (20-40) of having higher spending score around (40-100) and the age (60-70) of having balanced spending score around (40-60) respectively .

ii) Data Preprocessing

```
1 #converting gender to binary
2 segment.Gender=np.where(segment["Gender"]=="Male",1,0)
3 segment.Gender
```

	Gender
0	1
1	1
2	0
3	0
4	0
...	...
195	0
196	0
197	1
198	1
199	1

200 rows × 1 columns

```
1 segment.head()
```

	Gender	Age	AnnualIncome	SpendingScore
0	1	19	15	39
1	1	21	15	81
2	0	20	16	6
3	0	23	16	77
4	0	31	17	40

Next steps:

[Generate code with segment](#)
[View recommended plots](#)
[New interactive sheet](#)

```

1 # Scaling desired columns for modelling
2 scaler=StandardScaler()
3 scaled_val=scaler.fit_transform(segment[["AnnualIncome","SpendingScore"]])

```

```
1 scaled_val[:5]
```

```

array([[ -1.73899919, -0.43480148],
       [ -1.73899919,  1.19570407],
       [ -1.70082976, -1.71591298],
       [ -1.70082976,  1.04041783],
       [ -1.66266033, -0.39597992]])

```

```

1 # Creating a new dataframe with out Gender variable
2 features=pd.DataFrame(scaled_val,columns=segment.columns[2:4].tolist())
3 features

```

	AnnualIncome	SpendingScore
0	-1.738999	-0.434801
1	-1.738999	1.195704
2	-1.700830	-1.715913
3	-1.700830	1.040418
4	-1.662660	-0.395980
...
195	2.268791	1.118061
196	2.497807	-0.861839
197	2.497807	0.923953
198	2.917671	-1.250054
199	2.917671	1.273347

200 rows × 2 columns

Next steps: [Generate code with features](#) [View recommended plots](#) [New interactive sheet](#)

iii) Model Building & Evaluation

~: DBSCAN groups observations into clusters of high density ,which does not make use of k-clusters.

```

1 #DBSCAN Algo...
2 # eps is maximum distance b/w datapoints , min_samples for core datapoints.
3 dbscan=DBSCAN(eps=0.5,min_samples=10,metric="euclidean")
4 dbscan.fit(features)
5 ypre=dbscan.fit_predict(features)

```

```
1 dbscan
```

```

DBSCAN
DBSCAN(min_samples=10)

```

```
1 ypre
```

```

array([ -1,  0,  1,  0, -1,  0,  1, -1,  1,  0,  1, -1,  1,  0,  1,  0, -1,
         0, -1, -1, -1,  0,  1,  0,  1,  0,  0,  0,  0,  0,  1,  0,  1,  0,
         1,  0,  1,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
         0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
         0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
         0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
         0,  0,  0,  0,  2,  0,  2,  0,  2,  3,  2,  3,  2,  0,  2,  3,  2,
         3,  2,  3,  2,  3,  2,  0,  2,  3,  2,  0,  2,  3,  2,  3,  2,  3,
         2,  3,  2,  3,  2,  3,  2,  0,  2,  3,  2,  3,  2,  3,  2,  3,  2,
         3,  2,  3,  2,  3,  2,  3,  2,  3,  2, -1,  2,  3,  2, -1,  2,  3,
        -1,  3, -1,  3, -1, -1, -1, -1, -1, -1, -1, -1, -1])

```

~:Outliers are denoted with "-1"


```
1 dbscan.labels_
```

```

array([[-1,  0,  1,  0, -1,  0,  1, -1,  1,  0,  1, -1,  1,  0,  1,  0, -1,
        0, -1, -1, -1,  0,  1,  0,  1,  0,  0,  0,  0,  1,  0,  1,  0,
        1,  0,  1,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  2,  0,  2,  0,  2,  3,  2,  3,  2,  0,  2,  3,  2,
        3,  2,  3,  2,  3,  2,  0,  2,  3,  2,  0,  2,  3,  2,  3,  2,  3,
        2,  3,  2,  3,  2,  3,  2,  0,  2,  3,  2,  3,  2,  3,  2,  3,  2,
        3,  2,  3,  2,  3,  2,  3,  2,  3,  2, -1,  2,  3,  2, -1,  2,  3,
        -1,  3, -1,  3, -1, -1, -1, -1, -1, -1, -1, -1, -1])

1 # Plotting for clusters visualization
2 clr=["blue","green","red","orange","purple"]
3
4 pt.figure(figsize=(10,6))
5 for i in np.unique(ypre):
6     if i==-1:
7         pt.scatter(features[ypre== i]['AnnualIncome'],features[ypre== i]['SpendingScore'],s=50,c="black",label="Noise")
8     else:
9         pt.scatter(features[ypre== i]['AnnualIncome'],features[ypre== i]['SpendingScore'],label=f"Cluster {i}",c=clr[i])
10
11 pt.xlabel('AnnualIncome')
12 pt.ylabel('SpendingScore')
13 pt.title('DBSCAN Clustering')
14 pt.legend()
15 pt.show()

```



~: It's to be interpreted that black color datapoints represents outliers/noise in the clusters.

```

1 #Evaluation Metrics..
2 ss=metrics.silhouette_score(features,dbscan.labels_)
3 print("Silhouette_Score Coefficient : {:.2f}".format(ss))

```

```

Silhouette_Score Coefficient : 0.41

```

~:Silhouette_Score ranges from -1 to 1, which near to one is best and near to -1 is worst. Since we got coefficient as **0.41** in which datapoints are very Moderately compact with the clusters.

Double-click (or enter) to edit

Double-click (or enter) to edit

1 Start coding or [generate](#) with AI.